

USED CAR PRICE PREDICTION

A project report submitted to

DR. SUGANYA G

SCHOOL OF COMPUTER SCIENCE and ENGINEERING

is a partial fulfillment for the course

CSE4036 – MACHINE LEARNING



| Name | Reg. No. | Phone No. |
|-------------------|-----------|------------|
| Atul Patel | 20MIA1134 | 9179210782 |
| Shashank Pandey | 20MIA1147 | 9718630685 |
| Piyali Saha | 20MIA1066 | 6290612813 |
| Chetan Srivastava | 20MIA1122 | 6307696635 |

Abstract

New car prices in the industry are set by manufacturers and additional costs to the government in the form of taxes. Customers who purchase this new vehicle can therefore be confident that the amount they are investing is reasonable. However, sales of used cars increased worldwide as the prices of new cars skyrocketed and customers could no longer afford to buy them. Therefore, there is an urgent need for a used car price prediction system that effectively determines the value of a car based on various characteristics. This existing system involves a process in which sellers randomly set prices, and buyers know nothing about cars and their value in today's scenario. In fact, the seller also does not know the current value of the car or the price at which it should be sold. To solve this problem, we developed a model, which is very effective. – A car price prediction has been a high interest research area, as it requires noticeable effort and knowledge of the field expert. The model predicts used cars reasonable prices based on multiple aspects, including vehicle mileage, year of manufacturing, fuel consumption, transmission, fuel type, and engine size. This model can benefit sellers, buyers, and car manufacturers in the used cars market. Regression algorithms are used because they give continuous values as output rather than categorized values. This makes it possible to predict the actual price of a car rather than its price range. A user interface has also been developed that captures input from any user and displays the car price according to the user's input.

Key Words

- Linear Regression
- Decision Tree Regressor
- Random forest Regressor
- Gradient boosting
- XG boost regressor
- Hyper-parameter tuning
- Cross- validation

Introduction

Due to the many factors that determine the price of used cars on the market, it is difficult to determine whether listing prices for used cars is a difficult task. The focus of this project is to develop this machine learning model that can accurately predict used car prices based on used car characteristics and make informed purchases. We implement and evaluate different learning methods on a dataset consisting of retail prices of various makes and models. Compare the performance of different machine learning algorithms such as Linear Regression, Decision Tree Regressor, Random-forest, Gradient boosting, XG boost regressor and choose the best. Depending on various parameters, the price of the car is determined. Car price prediction is somehow interesting and popular problem. Accurate car price prediction involves expert knowledge, because price usually depends on many distinctive features and factors. Typically, most significant ones are brand, year, power and mileage. The fuel type used in the car as well as fuel consumption per mile highly affect price of a car due to a frequent change in the price of a fuel. In this project, we applied different methods and techniques in order to achieve higher precision of the used car price prediction.

Background study

Predicting price of a used cars has been studied extensively in various researches. Listian discussed, in her paper written for Master thesis, that regression model that was built using Support Vector Machines (SVM) can predict the price of a car that has been leased with better precision than multivariate regression or some simple multiple regression. This is on the grounds that Support Vector Machine (SVM) is better in dealing with datasets with more dimensions and it is less prone to overfitting and underfitting. The weakness of this research is that a change of simple regression with more advanced SVM regression was not shown in basic indicators like mean, variance or standard deviation.

Another approach was given by Richardson in his thesis work. His theory was that car producers produce more durable cars. Richardson applied multiple regression analysis and demonstrated that hybrid cars retain their value for longer time than traditional cars. This has roots in environmental concerns about the climate and it gives higher fuel efficiency. They look at how supervised machine learning techniques can be used to estimate the price of second-hand cars in Mauritius in this study. The forecasts are based on historical data taken from daily publications. To make the predictions, various techniques such as linear regression analysis, were employed.

According to author Sameerchand, car price estimates on historical data gathered from daily newspapers. For estimating the price of cars, they employed supervised machine learning algorithms. Other methods that have been employed include multiple linear regression, and

various decision tree algorithms. The best algorithm for prediction was identified after comparing all four algorithms. They had some issues comparing the algorithms, but they succeeded to do so.

Wu et al. in his study, exhibit automobile price prediction using a neural fuzzy knowledge-based system. They projected a model that has similar outcomes to the simple regression model by taking into account the following attributes: brand, year of manufacturing, and kind of engine. They have developed an expert system called ODAV (Optimal Distribution of Auction Automobiles) because there is a strong demand for car dealers to sell leased vehicles at the end of the lease year. This system gives insights into the best prices for vehicles, as well as the location where the best price can be gained. Regression model based on k-nearest neighbor machine learning algorithm was used to predict the price of a car. This system has a tendency to be exceptionally successful since more than two million vehicles were exchanged through it.

Noor and Jan build a model for car price prediction by using multiple linear regression. The dataset was created during the two-months period and included the following features: price, cubic capacity, exterior color, date when the ad was posted, number of ad views, power steering, mileage in kilometer, rims type, type of transmission, engine type, city, registered city, model, version, make and model year. After applying feature selection, the authors considered only engine type, price, model year and model as input features. With the given setup authors were able to achieve prediction accuracy of 98%. In the related work shown above, authors proposed prediction model based on the single machine learning algorithm.

Peerun et al. conducted study to assess the neural network's performance in predicting used automobile prices. However, especially on higher-priced cars, the estimated value is not very close to the real price. In forecasting the price of a used car, they found that support vector machine regression outperformed neural networks and linear regression by a little margin.

However, it is observed that single machine learning algorithm approach did not give remarkable prediction results and could be enhanced by assembling various machine learning methods in an ensemble to give more precision result.

Idea and Proposed Methodology

Ideology

Nearly 40% of India's Population falls under the middle-class category. - Survey by Krishnan and Hatekar. Nowadays the ridesharing market is functioning on resold cars and in today's world customized cars have become the next new thing.

Automobile industry is one of the markets where used cars are bought and sold, and in today's market Predicting the resale value of used cars is an important aspect. With the ever-growing number of used car sales thanks to the booming ridesharing market, it is of paramount importance that the buyer/seller's rupee is valued.

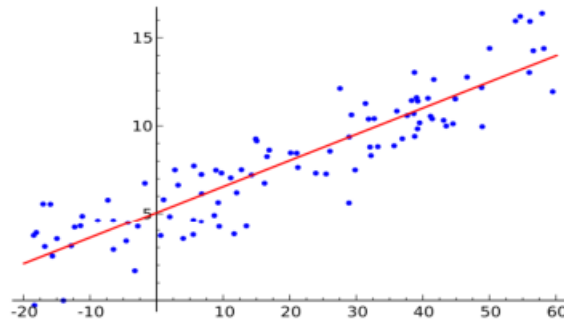
Methodology

There are two primary phases in the system:

1. Training phase: The system is trained by using the data in the data set and fits a model (line/curve) based on the algorithm chosen accordingly.
2. Testing phase: the system is provided with the inputs and is tested for its working. The accuracy is checked. And therefore, the data that is used to train the model or test it, has to be appropriate. The system is designed to detect and predict price of used car and hence appropriate algorithms must be used to do the two different tasks. Before the algorithms are selected for further use, different algorithms were compared for its accuracy. The well-suited one for the task was chosen.

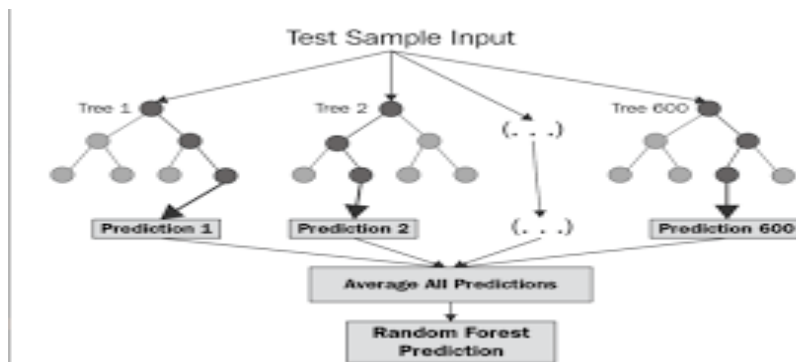
The main goal is to give users an accurate estimate of how much has to be paid for the given vehicle. The model may give the customer a record of possibilities for various automobiles based on the details of the automobile the customer wants. The system assists in providing the customer with sufficient data to help him to reach a conclusion. The used automobile market is expanding at an exponential rate, and vehicle vendors may profit from this by offering incorrect prices to capitalize on the demand. As a result, a system that can predict the price of a car based on its parameters while also taking into consideration the costs of competing vehicles is necessary. Our system fills in the gaps by providing buyers and sellers with an estimate of the car's value based on the best algorithm available for price prediction. This project aims to develop a good regression model to offer accurate prediction of car price. In order to do this, we need some previous data of used cars for which we use price and some other standard attributes. Car price is considered as the dependent variable while other attributes as the independent variables. Random Forest is an ensemble learning based regression model. It uses a model called decision tree, specifically as the name suggests, multiple decision trees to generate the ensemble model which collectively produces a prediction. The benefit of this model is that the trees are produced in parallel and are relatively uncorrelated, thus producing good results as each tree is not prone to individual errors of other trees. This uncorrelated behavior is partly ensured by the use of Bootstrap Aggregation or bagging providing the randomness required to produce robust and uncorrelated trees. This model was hence chosen to account for the large number of features in the dataset and compare a bagging technique with the following gradient boosting methods.

- LINEAR REGRESSION



Regression is a technique for predicting a dependent factor using independent factors. The approach is typically employed for predicting and calculating correlations between independent and dependent factors. The regression model establishes the relationship between independent factors and dependent factors i.e., linearly or exponentially. Linear regression is a sort of regression analysis in which there is only one independent variable and the independent(x) and dependent(y) variables can be bound in a linear relationship. In the graph above, the red line is called the best fit straight line. We aim to plot a line that best predicts the data-points based on the given data-points.

- RANDOM FOREST REGRESSION



It is a supervised learning algorithm that uses an ensemble learning method for classification and regression. Random forest is a bagging technique and not a boosting technique. The trees in random forests run in parallel, meaning is no interaction between these trees while building the trees. Random forest operates by constructing a multitude of decision trees at training time and outputting the class that's the mode of the classes (classification) or mean prediction (regression) of the individual trees.

It is one of the most accurate learning algorithms available. For many data sets, it produces a highly accurate classifier. It runs efficiently on large databases. It can handle thousands of input variables without variable deletion. It gives estimates of what variables are important in the classification. It generates an internal unbiased estimate of the generalization error as the forest building progresses.

It has an effective method for estimating missing data and maintains accuracy when a large proportion of the data are missing.

We will use the sklearn module for training our random forest regression model, specifically the RandomForestRegressor function using different parameters, few of those parameters are:

n_estimators — the number of decision trees you will be running in the model

criterion — this variable allows you to select the criterion (loss function) used to determine model outcomes. We can select from loss functions such as mean squared error (MSE) and mean absolute error (MAE). The default value is MSE.

max_depth — this sets the maximum possible depth of each tree

max_features — the maximum number of features the model will consider when determining a split

bootstrap — the default value for this is True, meaning the model follows bootstrapping principles

max_samples — This parameter assumes bootstrapping is set to True, if not, this parameter doesn't apply. In the case of True, this value sets the largest size of each sample for each tree.

- DECISION TREE REGRESSOR

A Decision Tree is a supervised learning algorithm. It is a graphical representation of all the possible solutions. All the decisions were made based on some conditions. Decision Trees are a type of

Supervised Machine Learning (that is you explain what the input is and what the corresponding output is in the training data) where the data is continuously split according to a certain parameter. The tree can be explained by two entities, namely decision nodes and leaves. The leaves are the decisions or the final outcomes. And the decision nodes are where the data is split.

There are two main types of Decision Trees:

1. **Classification trees** (Yes/No types)

Here, the outcome is a variable like 'fit' or 'unfit'. Here the decision variable is **Categorical**.

2. **Regression trees** (Continuous data types)

Here the decision or the outcome variable is **Continuous**. **Working**. There are many algorithms out there which construct Decision Trees, but one of the best is called as **ID3 Algorithm**. ID3 Stands for **Iterative Dichotomiser 3**

3. Entropy:

$$H(S) = \sum_{x \in X} p(x) \log_2 \frac{1}{p(x)}$$

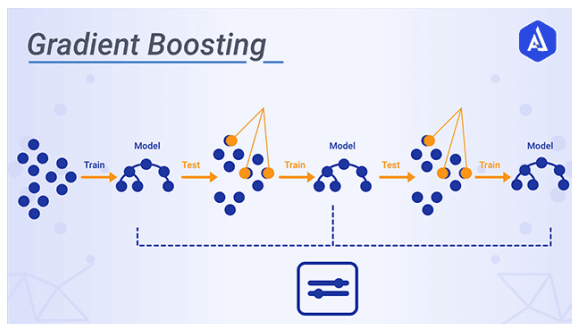
Entropy, also called as Shannon Entropy is denoted by $H(S)$ for a finite set S , is the measure of the amount of uncertainty or randomness in data. Intuitively, it tells us about the predictability of a certain event. Entropy to measure discriminatory power of an attribute for classification task. It defines the amount of randomness in attribute for classification task. Entropy is minimal means the attribute appears close to one class and has a good discriminatory power for classification

4. Information Gain:

$$IG(S, A) = H(S) - H(S, A)$$

Information gain is also called as Kullback-Leibler divergence denoted by $IG(S, A)$ for a set S is the effective change in entropy after deciding on a particular attribute A . It measures the relative change in entropy with respect to the independent variables.

- GRADIENT BOOSTING



Gradient boosting is a technique used in creating models for prediction. The technique is mostly used in regression and classification procedures. Prediction models are often presented as decision trees for choosing the best prediction. Gradient boosting presents model building in stages, just like other boosting methods, while allowing the generalization and optimization of differentiable loss functions. The gradient boosting is also known as the statistical prediction model. It works

quite similarly to other boosting methods even though it allows the generalization and optimization of the differential loss functions. One uses gradient boosting primarily in the procedures of regression and classification.

Gradient boosting creates prediction-based models in the form of a combination of weak prediction models. Weak hypotheses are parameters whose performance is slightly better than the randomly made choices. The gradient boosting method has witnessed many further developments to optimize the cost functions.

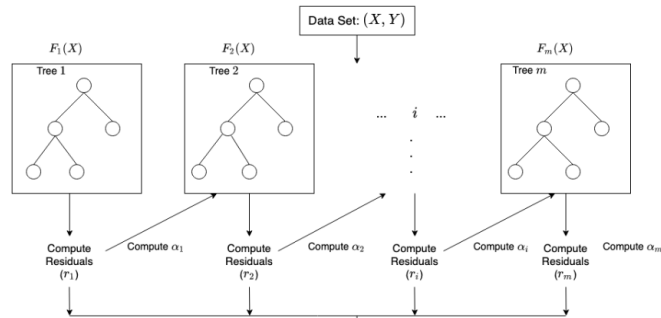
- XG BOOST

XG Boost stands for extreme Gradient Boosting. XG Boost is an extension to gradient boosted decision trees (GBM) and specially designed to improve speed and performance.

XG Boost is a popular and efficient open-source implementation of the gradient boosted trees algorithm. Gradient boosting is a supervised learning algorithm, which attempts to accurately predict a target variable by combining the estimates of a set of simpler, weaker models.

When using gradient boosting for regression, the weak learners are regression trees, and each regression tree maps an input data point to one of its leafs that contains a continuous score. XG Boost minimizes a regularized (L1 and L2) objective function that combines a convex loss function (based on the difference between the predicted and target outputs) and a penalty term for model complexity (in other words, the regression tree functions). The training proceeds iteratively, adding new trees that predict the residuals or errors of prior trees that are then combined with previous trees to make the final prediction. It's

called gradient boosting because it uses a gradient descent algorithm to minimize the loss when adding new models.



$$F_m(X) = F_{m-1}(X) + \alpha_m h_m(X, r_{m-1}).$$

where α_i and r_i are the regularization parameters and residuals computed with the i^{th} tree respectively, and h_i is a function that is trained to predict residuals, r_i using X for the i^{th} tree. To compute α_i we use the residuals computed, r_i and compute the following: $\arg \min_{\alpha} = \sum_{i=1}^m L(Y_i, F_{i-1}(X_i) + \alpha h_i(X_i, r_{i-1}))$ where $L(Y, F(X))$ is a differentiable loss function.

Requirements

- Hardware requirements
- Operating system- Windows 7,8,10
- Processor- dual core 2.4 GHz (i5 or i7 series Intel processor or equivalent AMD)
- RAM-4GB
- Software Requirements
- Python
- PyCharm
- PIP 2.7
- Jupyter Notebook
- Chrome

Implementation

jupyter used-car-prices-prediction (1) (1) Last Checkpoint: Last Sunday at 10:43 AM (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 C

```
In [72]: #Importing the necessary Libraries

import numpy as np
import pandas as pd
import seaborn as sb
import matplotlib.pyplot as plt
from sklearn.preprocessing import *
from sklearn.model_selection import *
from sklearn.ensemble import *
from sklearn.feature_selection import *
from sklearn.metrics import *
from sklearn.tree import *
from sklearn.svm import *
from sklearn.neighbors import *
from sklearn.linear_model import *
from sklearn.model_selection import KFold
import xgboost as xgb
from sklearn.ensemble import RandomForestRegressor, BaggingRegressor, AdaBoostRegressor, GradientBoostingRegressor

#To avoid unnecessary warning
import warnings
with warnings.catch_warnings():
    warnings.filterwarnings("ignore")
```

DATA GATHERING

loading the training dataset

```
In [73]: train = pd.read_csv('train-data.csv')
```

```
In [74]: train
```

Out[74]:

| | Unnamed: 0 | Name | Location | Year | Kilometers_Driven | Fuel_Type | Transmission | Owner_Type | Mileage | Engine | Power | Seats | New_Price | Price |
|---|------------|----------------------------------|----------|------|-------------------|-----------|--------------|------------|------------|---------|-----------|-------|-----------|-------|
| 0 | 0 | Maruti Wagon R LXI CNG | Mumbai | 2010 | 72000 | CNG | Manual | First | 26.6 km/kg | 998 CC | 58.16 bhp | 5.0 | NaN | 1.75 |
| 1 | 1 | Hyundai Creta 1.6 CRDI SX Option | Pune | 2015 | 41000 | Diesel | Manual | First | 19.67 kmpl | 1582 CC | 126.2 bhp | 5.0 | NaN | 12.50 |
| 2 | 2 | Honda Jazz V | Chennai | 2011 | 48000 | Petrol | Manual | First | 18.2 kmpl | 1199 CC | 88.7 bhp | 5.0 | 8.61 Lakh | 4.50 |
| 3 | 3 | Maruti Ertiga VDI | Chennai | 2012 | 87000 | Diesel | Manual | First | 20.77 kmpl | 1248 CC | 88.76 bhp | 7.0 | NaN | 6.00 |

- Here we have imported all the libraries and loaded the dataset here we can see the dataset

jupyter ML PROJECT Last Checkpoint: a minute ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 C

Run Markdow

| | | | | | | | | | | | | | | |
|------|------|-----------------------|-----------|------|-------|--------|--------|--------|------------|---------|----------|-----|-----|------|
| 6016 | 6016 | Mahindra Xylo D4 BSIV | Jaipur | 2012 | 55000 | Diesel | Manual | Second | 14.0 kmpl | 2498 CC | 112 bhp | 8.0 | NaN | 2.90 |
| 6017 | 6017 | Maruti Wagon R VXi | Kolkata | 2013 | 46000 | Petrol | Manual | First | 18.9 kmpl | 998 CC | 67.1 bhp | 5.0 | NaN | 2.85 |
| 6018 | 6018 | Chevrolet Beat Diesel | Hyderabad | 2011 | 47000 | Diesel | Manual | First | 25.44 kmpl | 936 CC | 57.6 bhp | 5.0 | NaN | 2.50 |

6019 rows x 14 columns

In [138]: train.shape

Out[138]: (6019, 14)

checking the dtype and null values

In [139]:

```
train = pd.read_csv('train-data.csv', index_col=0)
train = train.reindex(np.random.permutation(train.index))
print("TRAIN SHAPE: ", train.shape)
train.info()
train.head()
```

TRAIN SHAPE: (6019, 13)

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6019 entries, 625 to 5706
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Name                   6019 non-null  object
1   Location               6019 non-null  object
2   Year                   6019 non-null  int64
3   Kilometers_Driven      6019 non-null  int64
4   Fuel_Type              6019 non-null  object
5   Transmission           6019 non-null  object
6   Owner_Type             6019 non-null  object
7   Mileage                6017 non-null  object
8   Engine                 5983 non-null  object
9   Power                  5983 non-null  object
10  Seats                  5977 non-null  float64
11  New_Price              824 non-null   object
12  Price                  6019 non-null  float64
dtypes: float64(2), int64(2), object(9)
memory usage: 658.3+ KB
```

Out[139]:

| | Name | Location | Year | Kilometers_Driven | Fuel_Type | Transmission | Owner_Type | Mileage | Engine | Power | Seats | New_Price | Price |
|------|----------------------------|-----------|------|-------------------|-----------|--------------|------------|-----------|---------|-----------|-------|------------|-------|
| 625 | Honda City V MT | Pune | 2013 | 127000 | Petrol | Manual | Second | 16.8 kmpl | 1497 CC | 118.3 bhp | 5.0 | NaN | 4.50 |
| 2648 | Honda BRV i-VTEC V MT | Pune | 2016 | 9200 | Petrol | Manual | First | 15.4 kmpl | 1497 CC | 117.3 bhp | 7.0 | 13.58 Lakh | 8.00 |
| 3894 | Toyota Innova 2.5 GX 7 STR | Hyderabad | 2010 | 250000 | Diesel | Manual | Second | 12.8 kmpl | 2494 CC | 102 bhp | 7.0 | NaN | 6.50 |
| 4797 | Hyundai i20 1.2 Asta | Chennai | 2011 | 47000 | Petrol | Manual | First | 17.0 kmpl | 1197 CC | 80 bhp | 5.0 | NaN | 4.50 |

- We can observe that shape of our dataset is 6019 rows and 14 columns of which 13 columnn are independent and 1 which is a selling price is dependent or target column
- our dataset contain column such as location, year, kilometers driven, fuel type, Transmission , owner

Type, mileage, engine, power, seats, new_price, price .

- In our dataset year , kilometere driven , engine , mileage 4 coloumn are numerical 8 are categorical .
- In our dataset 2 values are float and 2 of them are integers and 9 of them are object

20PROJECT.ipynb

jupyter ML PROJECT Last Checkpoint: 6 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 C

Run

DATA CLEANING AND ASSESSMENT

Ensure to apply the changes to both the test and train sets inorder to maintain uniformity.

Quality issue 1: Missing values

We can see that there are missing values in the following columns:

- Engine
- Power
- Seats
- New_Price

percentage of missing values

```
In [141]: #percentage of missing values
percent_missing = train.isnull().sum() * 100 / len(train)
print(percent_missing)
```

| | |
|-------------------|-----------|
| Name | 0.000000 |
| Location | 0.000000 |
| Year | 0.000000 |
| Kilometers_Driven | 0.000000 |
| Fuel_Type | 0.000000 |
| Transmission | 0.000000 |
| Owner_Type | 0.000000 |
| Mileage | 0.033228 |
| Engine | 0.598106 |
| Power | 0.598106 |
| Seats | 0.697790 |
| New_Price | 86.310018 |
| Price | 0.000000 |

dtype: float64

Here we can see that train new price attributes , seats is having 86.31%missing values

```
In [142]: percent_missing = test.isnull().sum() * 100 / len(test)
print(percent_missing)
```

| | |
|-------------------|----------|
| Name | 0.000000 |
| Location | 0.000000 |
| Year | 0.000000 |
| Kilometers Driven | 0.000000 |

- Here we have cleaned the data, have found the missing values, we can observe the new price attribute contains 86% missing values hence we have dropped the column.
- Now we can see the the missing % is now 0.

Jupyter ML PROJECT Last Checkpoint: 6 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

Feature engineering

Feature engineering is the pre-processing step of machine learning, which is used to transform raw data into features that can be used for creating a predictive model using Machine learning

```
In [160]: x = pd.read_csv('trainfinal.csv')
          print(x.shape)
          x.head()
```

```
(5874, 13)
```

```
Out[160]:
```

| | Unnamed: 0 | Name | Location | Year | Kilometers_Driven | Fuel_Type | Transmission | Owner_Type | Mileage | Engine | Power | Seats | Price |
|---|------------|----------------------------|-----------|------|-------------------|-----------|--------------|------------|---------|--------|-------|-------|-------|
| 0 | 625 | Honda City V MT | Pune | 2013 | 127000 | Petrol | Manual | Second | 16.8 | 1497.0 | 116.3 | 5.0 | 4.50 |
| 1 | 2648 | Honda BRV i-VTEC V MT | Pune | 2016 | 9200 | Petrol | Manual | First | 15.4 | 1497.0 | 117.3 | 7.0 | 8.00 |
| 2 | 3894 | Toyota Innova 2.5 GX 7 STR | Hyderabad | 2010 | 250000 | Diesel | Manual | Second | 12.8 | 2494.0 | 102.0 | 7.0 | 6.50 |
| 3 | 4797 | Hyundai i20 1.2 Asta | Chennai | 2011 | 47000 | Petrol | Manual | First | 17.0 | 1197.0 | 80.0 | 5.0 | 4.50 |
| 4 | 3813 | Hyundai Grand i10 Sportz | Mumbai | 2013 | 26000 | Petrol | Manual | First | 18.9 | 1197.0 | 82.0 | 5.0 | 3.75 |

dropping the unnamed:0 column

```
In [161]: x.drop(columns=['Unnamed: 0'],axis=1,inplace = True)
```

The name column has a diverse number of values. Let's break it down and extract the brand name of the car.

```
In [162]: x["breakdown"] = x.Name.str.split(" ")
          x["breakdown"].head()
```

```
Out[162]:
```

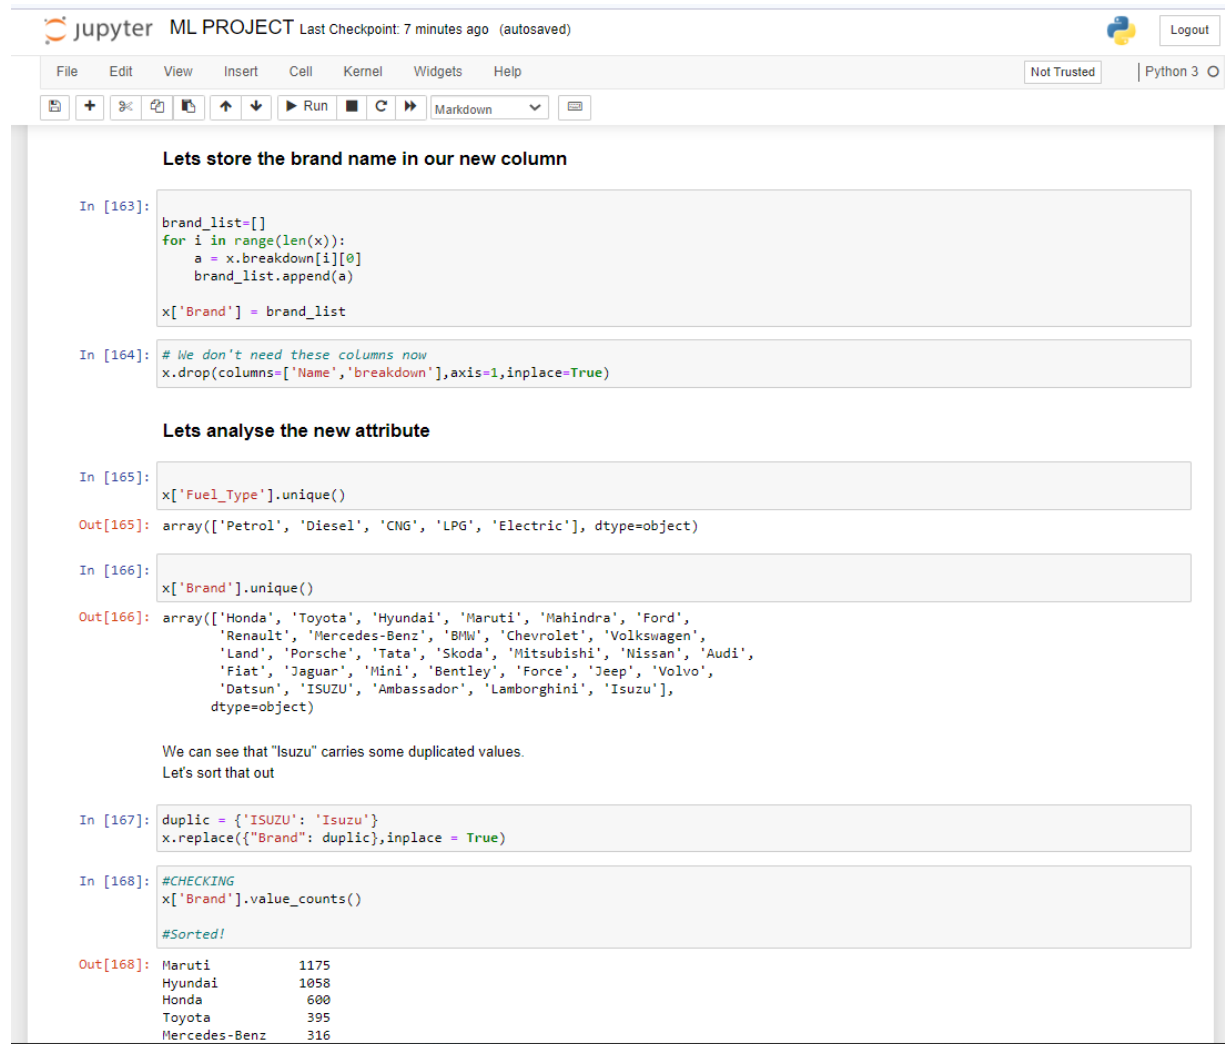
```
0      [Honda, City, V, MT]
1      [Honda, BRV, i-VTEC, V, MT]
2      [Toyota, Innova, 2.5, GX, 7, STR]
3      [Hyundai, i20, 1.2, Asta]
4      [Hyundai, Grand, i10, Sportz]
Name: breakdown, dtype: object
```

Lets store the brand name in our new column

- Here we have done the feature engineering with the preprocessed dataset and made the dataset fit for the modeling
- Here we have a breakdown of the name attribute in the comapany name and brand name

- And also analyse the new attributes such as brand and , fuel type etc

fu



The image shows a Jupyter Notebook interface with the title "ML PROJECT" and a status bar indicating "Last Checkpoint: 7 minutes ago (autosaved)". The interface includes a menu bar (File, Edit, View, Insert, Cell, Kernel, Widgets, Help) and a toolbar with icons for file operations, running, and saving. The notebook content is as follows:

Lets store the brand name in our new column

```
In [163]: brand_list=[]
          for i in range(len(x)):
              a = x.breakdown[i][0]
              brand_list.append(a)

          x['Brand'] = brand_list
```

```
In [164]: # We don't need these columns now
          x.drop(columns=['Name', 'breakdown'],axis=1,inplace=True)
```

Lets analyse the new attribute

```
In [165]: x['Fuel_Type'].unique()

Out[165]: array(['Petrol', 'Diesel', 'CNG', 'LPG', 'Electric'], dtype=object)
```

```
In [166]: x['Brand'].unique()

Out[166]: array(['Honda', 'Toyota', 'Hyundai', 'Maruti', 'Mahindra', 'Ford',
                'Renault', 'Mercedes-Benz', 'BMW', 'Chevrolet', 'Volkswagen',
                'Land', 'Porsche', 'Tata', 'Skoda', 'Mitsubishi', 'Nissan', 'Audi',
                'Fiat', 'Jaguar', 'Mini', 'Bentley', 'Force', 'Jeep', 'Volvo',
                'Datsun', 'ISUZU', 'Ambassador', 'Lamborghini', 'Isuzu'],
                dtype=object)
```

We can see that "Isuzu" carries some duplicated values.
Let's sort that out

```
In [167]: duplic = {'ISUZU': 'Isuzu'}
          x.replace({"Brand": duplic},inplace = True)
```

```
In [168]: #CHECKING
          x['Brand'].value_counts()

          #Sorted!
```

```
Out[168]: Maruti      1175
          Hyundai    1058
          Honda       600
          Toyota      395
          Mercedes-Benz 316
```

FEATURE IMPORTANCE RELATED TO TARGET

```
In [0]: # Find most important features relative to target Price
print("Find most important features relative to target")
corr = cars.corr()
corr.sort_values(["Price"], ascending = False, inplace = True)
print(corr.Price)

"""
Between Year and Price there is a positive correlation, meaning that the higher
is the Year (more recent), the higher is the Price (more recent cars cost more).
Between Price and Mileage there is a negative correlation, meaning that higher
is the mileage, lower is the Price (cars with high mileage cost less).

We obviously will get rid of Id attribute.
"""

Find most important features relative to target
Price      1.000000
Year       0.384824
Id        -0.022876
Mileage    -0.405779
Name: Price, dtype: float64
```

- Here we have to find the correlation between the independent and dependent variable
- Also we have done feature selection, like what is the features which are mainly important for our prediction model
- We can observe in the below feature selection that transmission is the highest priority factor followed by kilometer driven than fuel type.

Jupyter ML PROJECT Last Checkpoint: 7 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted | Python 3 C

4 9 15 26000 4 1 0 18.9 1197.0 82.0 5.0 3.75 10

feature selection

```
In [171]:
X_fs = x[['Location', 'Year', 'Kilometers_Driven', 'Fuel_Type', 'Transmission', 'Owner_Type',
         'Mileage', 'Engine', 'Power', 'Seats', 'Brand']]

y_fs = x['Price']

y_fs = y_fs*100
y_fs = y_fs.astype(int)

bestfeatures = SelectKBest(score_func=chi2, k=5)
fit = bestfeatures.fit(X_fs,y_fs)
dfscores = pd.DataFrame(fit.scores_)
dfcolumns = pd.DataFrame(X_fs.columns)
featureScores = pd.concat([dfcolumns,dfscores],axis=1)
featureScores.columns = ['Specs','Score'] #naming the dataframe columns
fea = pd.DataFrame(featureScores.nlargest(10,'Score'))
print(featureScores.nlargest(10,'Score')) #print 10 best features
```

| | Specs | Score |
|----|-------------------|--------------|
| 2 | Kilometers_Driven | 2.646175e+08 |
| 7 | Engine | 8.785410e+05 |
| 8 | Power | 1.170354e+05 |
| 10 | Brand | 6.291769e+03 |
| 5 | Owner_Type | 2.423341e+03 |
| 6 | Mileage | 2.241475e+03 |
| 1 | Year | 2.015923e+03 |
| 3 | Fuel_Type | 2.011687e+03 |
| 0 | Location | 1.808910e+03 |
| 4 | Transmission | 1.002323e+03 |

In []:

FEATURE IMPORTANCE RELATED TO TARGET

```
In [0]:
# Find most important features relative to target Price
print("Find most important features relative to target")
corr = cars.corr()
corr.sort_values(["Price"], ascending = False, inplace = True)
print(corr.Price)

"""
Between Year and Price there is a positive correlation, meaning that the higher
is the Year (more recent), the higher is the Price (more recent cars cost more).
Between Price and Mileage there is a negative correlation, meaning that higher
is the mileage, lower is the Price (cars with high mileage cost less).

We obviously will get rid of Id attribute.
"""

Find most important features relative to target
Price      1.000000
Year       0.384824
Id        -0.022876
Mileage    -0.405779
Name: Price, dtype: float64
```

- Here we have come to the model building part where we have split the data set into train and test with the ratio of 80:20.

- Then we have fitted the data by scaler transform
- Here we have applied various models which are as follows:-
 1. Linear regression & their metrics
 2. Decision
 3. Random
 4. gradient boosting
 5. xg boost

The image shows a Jupyter Notebook interface for an ML project. The notebook is titled 'ML PROJECT' and shows the following steps:

MODEL BUILDING

Preparing Training set

```
In [174]: X = np.array(x.drop(['Price'],axis = 1))
          Y = x.Price.values
```

splitting in train test

```
In [175]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.20, random_state=25)

print("X_test shape:", X_test.shape)
print("X_train shape:", X_train.shape)
print("Y_test shape: ", Y_test.shape)
print("Y_train shape:", Y_train.shape)

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

X_test shape: (1175, 11)
X_train shape: (4699, 11)
Y_test shape: (1175,)
Y_train shape: (4699,)
```

1. linear regression

```
In [176]: model = LinearRegression()
          model.fit(X_train, Y_train)
          pred = model.predict(X_test)

In [177]: model.score(X_test,pred)

Out[177]: 1.0

In [178]: import sklearn.metrics as metrics
```

LINEAR REGRESSION:-

Here we can see the model got fitted and we have metrics such as

1. MAE SCORE:- 3.77
2. MSE SCORE:- 37.58
3. R2 SCORE :- 0.58
4. RMSE:- 6.130

Here we can infer that the built model has produced the above metrics and this score such as MAE, MSE, and RMSE should be reduced ie on the other hand r2 score should be increased .

1. linear regression

```
In [176]: model = LinearRegression()
          model.fit(X_train, Y_train)
          pred = model.predict(X_test)
```

```
In [177]: model.score(X_test, pred)
```

Out[177]: 1.0

```
In [178]: import sklearn.metrics as metrics
```

here we can see that the metrices for the linear regression model is having mean absolute error, mean square error and r² score as follows :-

- MAE: 3.7705582986157813
- MSE: 37.58518408727565
- R2 score: 0.5801078926580725
- RMSE : 6.130676

from the above metrices we can conclude that model need to have be more optimal resulting low value of metrices

```
In [179]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

print("MAE: ", (metrics.mean_absolute_error(pred, Y_test)))
print("MSE: ", (metrics.mean_squared_error(pred, Y_test)))
print("R2 score: ", (metrics.r2_score(pred, Y_test)))
rmse = np.sqrt(metrics.mean_squared_error(Y_test, pred))
print("RMSE: % f" %(rmse))
```

MAE: 3.7705582986157813
MSE: 37.58518408727565
R2 score: 0.5801078926580725
RMSE : 6.130676

RMSE :commonly used measures for evaluating the quality of predictions. It shows how far predictions fall from measured true values If the RMSE for the test is much higher than that of the training set, it is likely that we have badly over fit the data, i.e. we have created a model that tests well in sample, but has little predictive value when tested out of sample.

DECISION TREE REGRESSOR :-

Here we can see the model got fitted and we have metrics such as

5. MAE SCORE:- 2.01
6. MSE SCORE:- 28.13

7. R2 SCORE :- 0.77

8. RMSE:- 5.30

Here also we can infer that the built model has produced the above metrics and this score such as MAE, MSE, and RMSE is reducing ie on the other hand r2 score should be increasing let's try some other model to reduce these metrics and increase the r2 score .

2. Decision Tree Regressor

```
In [180]: from sklearn.tree import DecisionTreeRegressor
```

```
model = DecisionTreeRegressor()  
model.fit(X_train, Y_train)  
pred = model.predict(X_test)
```

here we can observe that decision tree regressor is producing more accurate result as compared to linear model as metrics values of mean absolute error , mean square error values got reduced , and r2 score got increased

```
In [181]: from sklearn.metrics import mean_absolute_error  
print("MAE: ", (metrics.mean_absolute_error(pred, Y_test)))  
print("MSE: ", (metrics.mean_squared_error(pred, Y_test)))  
print("R2 score: ", (metrics.r2_score(pred, Y_test)))  
rmse = np.sqrt(metrics.mean_squared_error(Y_test, pred))  
print("RMSE : % f" %(rmse))
```

```
MAE:  2.019160283687943  
MSE:  28.134410009456264  
R2 score:  0.7781018946743732  
RMSE :  5.304188
```

Lower values of RMSE indicate better fit. The lower the RMSE, the better a given model is able to "fit" a dataset.

3. Random Forest Regressor

```
In [182]: from sklearn.ensemble import RandomForestRegressor  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import mean_absolute_error as mae  
  
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=.25, random_state=0)  
model = RandomForestRegressor(random_state=1)  
model.fit(X_train, Y_train)  
pred = model.predict(X_test)
```

```
In [183]: from sklearn.metrics import mean_absolute_error  
print("MAE: ", (metrics.mean_absolute_error(pred, Y_test)))  
print("MSE: ", (metrics.mean_squared_error(pred, Y_test)))  
print("R2 score: ", (metrics.r2_score(pred, Y_test)))  
rmse = np.sqrt(metrics.mean_squared_error(Y_test, pred))  
print("RMSE : % f" %(rmse))
```

```
MAE:  1.5017894787513368  
MSE:  10.85872853452442  
R2 score:  0.8942806738949218  
RMSE :  3.295258
```

Here, R2 score got increased, the higher the R-squared, the better the model fits your data

RANDOM FOREST REGRESSOR:-

Here we can see the model got fitted and we have metrics such as

- 9. MAE SCORE:- 1.50
- 10.MSE SCORE:- 10.85
- 11.R2 SCORE:- 0.89
- 12.RMSE:- 3.29

Here also we can infer that the built model has produced the above metrics and this score such as MAE, MSE, and RMSE is reducing ie on the other hand r2 score should be increasing let's try some other model to reduce these metrics and increase the r2 score .

4. Gradient Boosting Regressor

```
In [184]: model = GradientBoostingRegressor()
          model.fit(X_train, Y_train)
          pred = model.predict(X_test)

In [185]: from sklearn.metrics import mean_absolute_error
          print("MAE: ", (metrics.mean_absolute_error(pred, Y_test)))
          print("MSE: ", (metrics.mean_squared_error(pred, Y_test)))
          print("R2 score: ", (metrics.r2_score(pred, Y_test)))
          rmse = np.sqrt(metrics.mean_squared_error(Y_test, pred))
          print("RMSE : % f" %(rmse))

          MAE: 1.705947689373957
          MSE: 10.90467045731623
          R2 score: 0.8947487714525375
          RMSE : 3.302222

In [186]: import xgboost as xg
          from sklearn.metrics import mean_squared_error as MSE
          import warnings
          warnings.filterwarnings("ignore", category=DeprecationWarning)
```

5. xgb regressor

```
In [187]: xgb_r = xg.XGBRegressor(objective ='reg:linear',
          n_estimators = 100, seed = 123)

          # Fitting the model
          xgb_r.fit(X_train, Y_train)

          # Predict the model
          pred = xgb_r.predict(X_test)

          from sklearn.metrics import mean_absolute_error
          print("MAE: ", (metrics.mean_absolute_error(pred, Y_test)))
          print("MSE: ", (metrics.mean_squared_error(pred, Y_test)))
          print("R2 score: ", (metrics.r2_score(pred, Y_test)))
          # RMSE Computation
          rmse = np.sqrt(MSE(Y_test, pred))
          print("RMSE : % f" %(rmse))

          MAE: 1.368200062168464
          MSE: 9.43618017020281
          R2 score: 0.9115482386132322
          RMSE : 3.071837
```

GRADIENT BOOSTING REGRESSOR:-

Here we can see the model got fitted and we have metrics such as

13.MAE SCORE:- 1.50

14.MSE SCORE:- 10.90

15.R2 SCORE:- 0.84

16.RMSE:- 3.30

Here also we can infer that the built model has produced the above metrics and this score such as MAE, MSE, and RMSE is reducing ie on the other hand r2 score should be increasing let's try the model to reduce the metrics and increase the r2 score

XG BOOST REGRESSOR:-

Here we can see the model got fitted and we have metrics such as

17.MAE SCORE:- 1.36

18.MSE SCORE:- 9.43

19.R2 SCORE:- 0.91

20.RMSE:- 3.29

Here also we can infer that the built model has produced the above metrics and these scores such as MAE, MSE, and RMSE reached to the minimum value as compared to ie on the other hand r2 score is also high as compared to all let's try some other model to reduce these metrics and increase the r2 score

Lower values of RMSE indicate better fit. The lower the RMSE, the better a given model is able to "fit" a dataset. here in XG BOOST regressor it implies highest R2score and lowest RMSE among all other models

```
In [188]: #xgb_r.fit(X,Y)
```

```
In [189]: #from joblib import Parallel, delayed
#import joblib

# Save the model as a pickle in a file
#joblib.dump(xgb_r, 'model.pkl')

# Load the model from the file
```

AND THERE WE HAVE A WINNER TO GO THROUGH GRIDSEARCH FOR HYPER-PARAMETER TUNING!!

Tuning HyperParameters for xgboost

GRID SEARCH

Grid Search can be thought of as an exhaustive and computationally expensive method for selecting a model.

For example :

Searching 20 different parameter values for each of 4 parameters will require 160,000 trials of cross-validation. This equates to 1,600,000 model fits and 1,600,000 predictions if 10-fold cross validation is used

```
In [190]: #We use GridSearch for fine tuning Hyper Parameters

from sklearn.model_selection import *

n_estimator_val = np.arange(100,400,100).astype(int)
max_depth_val = [2,3,4]

grid_params = { 'learning_rate' : [0.1],
                 'n_jobs': [-1],
                 'n_estimators' : n_estimator_val,
                 'max_depth' : max_depth_val
               }
```

- Here we have done hyper parameter tuning using grid search and random random method
- We have found the best parameter such as learning rate =0.1, max depth =4, no. of estimator =300, no_jobs=1.
- And same parameter we have observed from random search algorithm which is mentioned below

To Display the Best Score

```
In [192]: gs_results.best_score_
```

```
Out[192]: 0.9071962570393127
```

To Display the Best Estimator

```
In [193]: gs_results.best_estimator_
```

```
Out[193]: XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None,
                        colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
                        early_stopping_rounds=None, enable_categorical=False,
                        eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',
                        importance_type=None, interaction_constraints='',
                        learning_rate=0.1, max_bin=256, max_cat_to_onehot=4,
                        max_delta_step=0, max_depth=4, max_leaves=0, min_child_weight=1,
                        missing=nan, monotone_constraints=('',), n_estimators=300,
                        n_jobs=-1, num_parallel_tree=1, predictor='auto', random_state=0,
                        reg_alpha=0, reg_lambda=1, ...)
```

To Display the Best Parameters

```
In [194]: gs_results.best_params_
```

```
Out[194]: {'learning_rate': 0.1, 'max_depth': 4, 'n_estimators': 300, 'n_jobs': -1}
```

We could settle here but let's try randomized search cross validation.

RANDOMIZED SEARCH

Randomized Search sets up a grid of hyperparameter values and selects random combinations to train the model and score. This allows you to explicitly control the number of parameter combinations that are attempted. The number of search iterations is set based on time or resources.

While it's possible that RandomizedSearchCV will not find as accurate of a result as GridSearchCV, it surprisingly picks the best result more often than not and in a fraction of the time it takes GridSearchCV would have taken. Given the same resources, Randomized Search can even outperform Grid Search.

```
In [195]: xgbr=xgb.XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=0.8, gamma=0.5,
importance_type='gain', learning_rate=0.1, max_delta_step=0,
max_depth=4, min_child_weight=10, missing=None, n_estimators=1900,
n_jobs=-1, nthread=None, objective='reg:linear', random_state=0,
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
silent=None, subsample=1.0, verbosity=1)
```

```
In [196]: folds = 4
param_comb = 10

skf = StratifiedKFold(n_splits=folds, shuffle = True, random_state = 1001)

params = { 'n_jobs': [-1],
'n_estimators' : n_estimator_val,
'learning_rate' : [0.1],
'min_child_weight': [9],
'gamma': [0.5],
'subsample': [0.6],
'colsample_bytree': [0.8, 1.0],
'max_depth': [3, 4]
}

xgb_regrsv = xgb.XGBRegressor()
```

```
In [197]: random_search = RandomizedSearchCV(xgb_regrsv, params, n_iter=param_comb, scoring='r2',
n_jobs=-1, cv=10 )
```

```
In [198]: random_search.fit(X_train, Y_train);
```

```
In [199]: random_search.best_score_
```

```
Out[199]: 0.9117690415341615
```

```
In [200]: random_search.best_estimator_
```

```
Out[200]: XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None,
colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1.0,
early_stopping_rounds=None, enable_categorical=False,
```

- After finding the best parameter we tuned it according to that parameter to get the best accuracy
- And we have also drawn the matrices and accuracy on the training set and testing set of the data
 1. Training set score =97%
 2. Testing set score=91%
- Metrics
 3. MAE SCORE:- 1.36
 4. MSE SCORE:- 9.15
 5. R2 SCORE:- 0.91
 6. RMSE:- 3.02

These is the minimum metrics that we have drawn among all model and with the xg boost also , before tuning the r2 score was less from 0.91 and paralley other score like MAE,MSE,RMSE we also little high

```
In [200]: random_search.best_estimator_
```

```
Out[200]: XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None,
                      colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1.0,
                      early_stopping_rounds=None, enable_categorical=False,
                      eval_metric=None, gamma=0.5, gpu_id=-1, grow_policy='depthwise',
                      importance_type=None, interaction_constraints='',
                      learning_rate=0.1, max_bin=256, max_cat_to_onehot=4,
                      max_delta_step=0, max_depth=4, max_leaves=0, min_child_weight=0,
                      missing=nan, monotone_constraints=('',), n_estimators=300,
                      n_jobs=-1, num_parallel_tree=1, predictor='auto', random_state=0,
                      reg_alpha=0, reg_lambda=1, ...)
```

After some more time of trying various hyperparameters and tuning.

```
In [201]: xgb_tuned = xgb.XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                      colsample_bynode=1, colsample_bytree=0.8, gamma=0.5,
                      importance_type='gain', learning_rate=0.1, max_delta_step=0,
                      max_depth=4, min_child_weight=9, missing=np.nan, n_estimators=300,
                      n_jobs=-1, nthread=None, objective='reg:linear', random_state=0,
                      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                      silent=True, subsample=1.0, verbosity=0)
```

```
In [202]: xgb_tuned.fit(X_train,Y_train)
y_pred =xgb_tuned.predict(X_test)

print("Training set accuracy: ",xgb_tuned.score(X_train,Y_train))
print("Test set accuracy   : ",xgb_tuned.score(X_test,Y_test))
```

```
Training set accuracy:  0.9766318073352496
Test set accuracy      :  0.9180405345235487
```

Now, lets check how various metrics have evaluated our model on the test set

```
In [203]: print("\t\tError Table")
print('Mean Absolute Error      : ', mean_absolute_error(Y_test, y_pred))
print('Mean Squared Error      : ', mean_squared_error(Y_test, y_pred))
print('Root Mean Squared Error : ', np.sqrt(mean_squared_error(Y_test, y_pred)))
print('R Squared Error         : ', r2_score(Y_test, y_pred))
```

```
Error Table
Mean Absolute Error      :  1.360548466045975
Mean Squared Error      :  9.153331821022785
Root Mean Squared Error :  3.0254473753517486
R Squared Error         :  0.9180405345235487
```

- **CROSS-VALIDATION:-**

After doing Hyper paramet tuning we have applied the cross validation on our and got the accuracy score of 84.26%

cross validation strategy

Cross Validation is used to assess the predictive performance of the models and to judge how they perform outside the sample to a new data set also known as test data. The motivation to use cross validation techniques is that when we fit a model, we are fitting it to a training dataset. Without cross validation we only have information on how does our model perform to our in-sample data. Ideally we would like to see how does the model perform when we have a new data in terms of accuracy of its predictions. In science, theories are judged by its predictive performance.

We use the `cross_val_score` function of Sklearn. However this function has not a shuffle attribut, we add then one line of code, in order to shuffle the dataset prior to cross-validation

```
In [208]: #Validation function
n_folds = 5

def rmsle_cv(model):
    kf = KFold(n_folds, shuffle=True, random_state=42).get_n_splits(X_train.values)
    rmse= np.sqrt(-cross_val_score(model, X_train.values, y_train.values, scoring="neg_mean_squared_error", cv = kf))
    return(rmse)
```

```
In [209]: from sklearn.model_selection import KFold, cross_val_score, train_test_split
random_model = RandomForestRegressor(n_estimators=300, max_depth=4)
#Fit
random_model.fit(X_train, Y_train)
xgb_accuracy = round(random_model.score(X_train,Y_train)*100,2)
print(round(xgb_accuracy,2),'%')
```

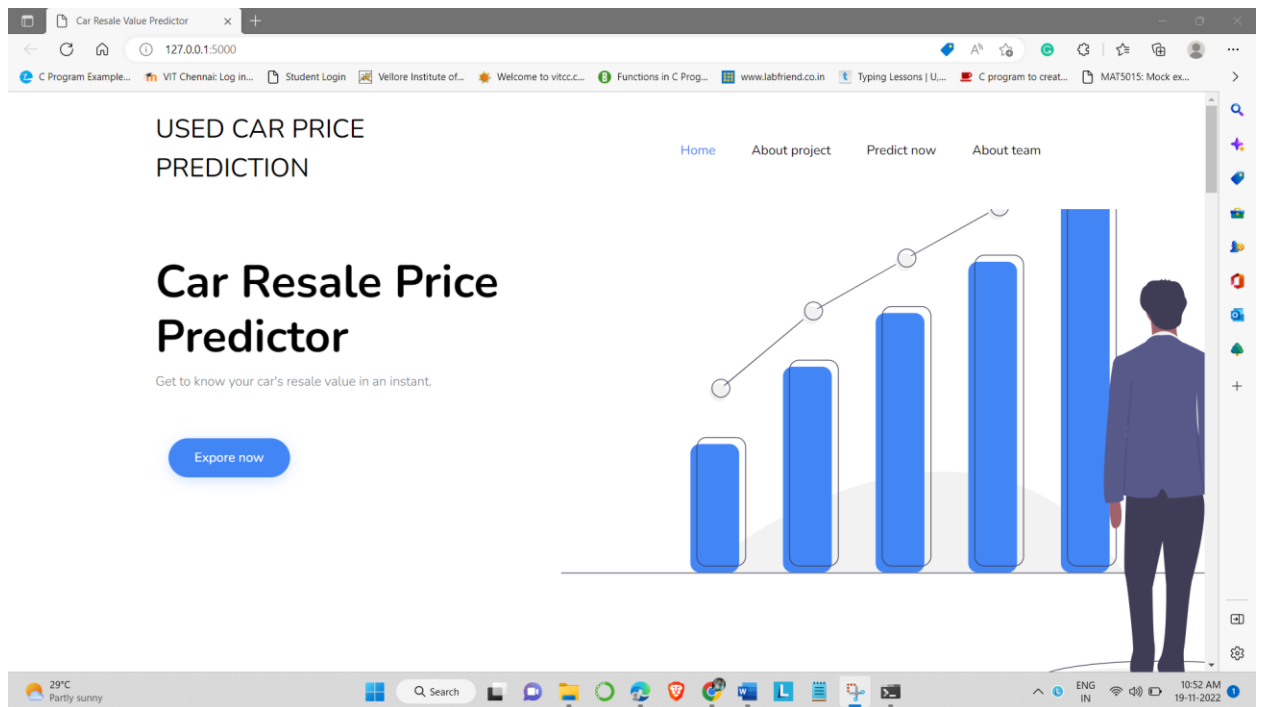
84.26 %

CONCLUSION

This model is ready to be deployed in a pipeline.

We have accomplished the task of building a good model for our used cars' prediction purposes.

- **Our Website :-**



Car Resale Value Predictor

127.0.0.1:5000


C Program Example... VIT Chennai Log in... Student Login Vellore Institute of... Welcome to vitcc... Functions in C Prog... www.labfriend.co.in Typing Lessons | U... C program to creat... MAT5015: Mock ex...

USED CAR PRICE PREDICTION

Home About project Predict now About team


About our project

Ideation behind our project




Introduction

Vehicle price prediction especially when the vehicle is used and not coming direct from the factory, is both a critical and important task. With increase in demand for used cars more and more vehicle buyers are finding alternatives of buying new cars. There is a need of accurate price prediction mechanism for the used cars. Prediction techniques of machine learning can be helpful in this regard. It is common to lease a car in many countries



Objective

The objective of the project is to find the possible resale value of a car based on its gearbox, model, brand, vehicle type, fuel type, mileage and engine power.



Abstract

The price of a new car in the industry is fixed by the manufacturer with some additional costs incurred by the Government in the form of taxes. So, customers buying a new car can be assured of the money they invest to be worthwhile. But, due to the increased prices of new cars and the financial incapability of the customers to buy them, Used Car sales are on a global increase. Therefore, there is an urgent

29°C Partly sunny

Q Search

ENG IN 10:55 AM 19-11-2022


Car Resale Value Predictor

127.0.0.1:5000

C Program Example... VIT Chennai Log in... Student Login Vellore Institute of... Welcome to vitcc... Functions in C Prog... www.labfriend.co.in Typing Lessons | U... C program to creat... MAT5015: Mock ex...


USED CAR PRICE PREDICTION

Home About project Predict now About team





Working Process

A car price prediction has been a high-interest research area, as it requires noticeable effort and knowledge of the field expert. A considerable number of distinct attributes are examined for reliable and accurate prediction. To build a model for predicting the price of used cars the applied machine learning techniques. Respective performances of different algorithms were then compared to find one that best suits the available data set. The final prediction model will be integrated into Python based Flask FrameWork. Furthermore, the model was evaluated using test data.



Methodolgy





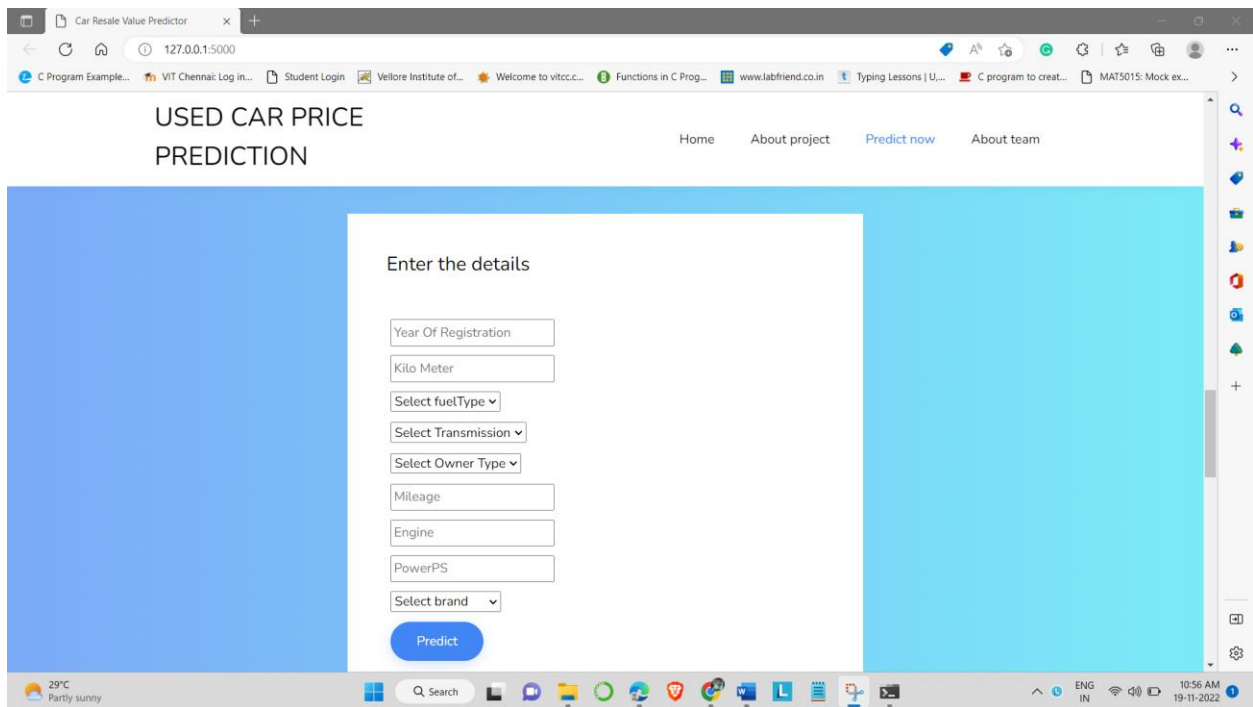
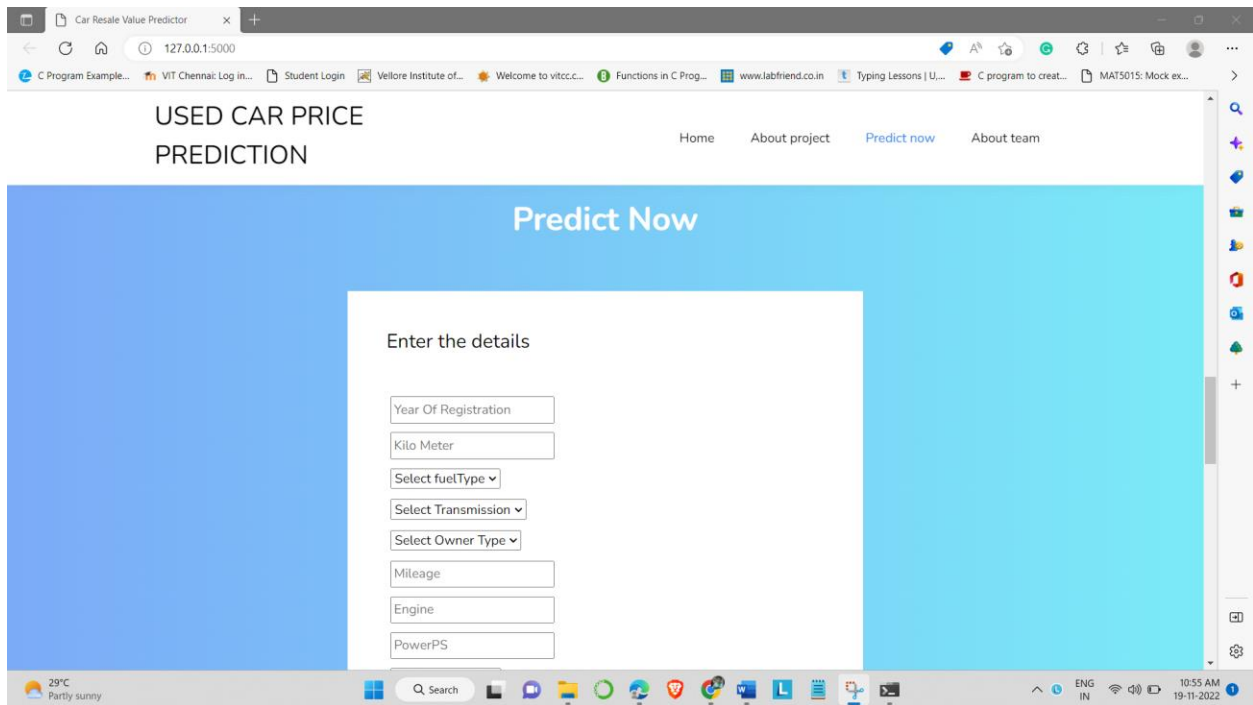
Future Scope

In future this machine learning model may bind with various website which can provide real time data for price prediction. Also we may add large historical data of car price which can help to improve accuracy of the machine learning model. We can build an android app as user interface for interacting with user. The proposed system will help to determine the accurate price of used car price prediction.

29°C Partly sunny

Q Search

ENG IN 10:55 AM 19-11-2022



PREDICTION:-

Car Resale Value Predictor

127.0.0.1:5000/y_predict

USED CAR PRICE PREDICTION

Home About project **Predict now** About team

Enter the details

2012

20000

Diesel

Manual

Second

16

800

116

Honda

Predict

The expected resale value is Rs. 4.56 lakh

29°C Partly sunny

Q Search

ENG IN 10:57 AM 19-11-2022


Car Resale Value Predictor

127.0.0.1:5000/y_predict


USED CAR PRICE PREDICTION

Home About project Predict now About team


Our Team




Amit Patel
INTERGRATED MTech CSE STUDENT AT VIT CHENNAI



Priyati Saha
INTERGRATED MTech CSE STUDENT AT VIT CHENNAI



Shashank Pandey
INTERGRATED MTech CSE STUDENT AT VIT CHENNAI



Dheeraj Srinivasan
INTERGRATED MTech CSE STUDENT AT VIT CHENNAI

29°C Partly sunny

Q Search

ENG IN 11:01 AM 19-11-2022

Results and Discussion

Results showing testing, metrics taken for checking and a complete analysis on the results obtained.

From the above

- **Linear Regression Model: -**

here we can see that the metrics for the linear regression model is having mean absolute error, mean square error and r² score as follows :-

- MAE: 1.503625622243836
- MSE: 11.342181490343576
- R2 score: 0.894980757373783
- RMSE : 3.367816

from the above metrics we can conclude that model need to have be more optimal resulting low value of metrics

```
[179]
... MAE: 3.7705582986157813
    MSE: 37.58518408727565
    R2 score: 0.5801078926580725
    RMSE : 6.130676
```

RMSE :commonly used measures for evaluating the quality of predictions. It shows how far predictions fall from measured true values If the RMSE for the test is much higher than that of the training set, it is likely that we have badly over fit the data, i.e. we have created a model that tests well in sample, but has little predictive value when tested out of sample.

- **Decision Tree Regressor: -**

```
... MAE: 2.019160283687943
    MSE: 28.134410009456264
    R2 score: 0.7781018946743732
    RMSE : 5.304188
```

Lower values of RMSE indicate better fit. The lower the RMSE, the better a given model is able to "fit" a dataset.

here we can observe that decision tree regressor is producing more accurate result as compared to linear model as metrics values of mean

absolute error, mean square error values got reduced , and r^2 score got increased

- **Radom Tree Regressor: -**

```
... MAE: 1.5017894787513368
    MSE: 10.85872853452442
    R2 score: 0.8942806738949218
    RMSE : 3.295258
```

Here, R^2 score got increased, the higher the R-squared, the better the model fits your data

- **Gradient Boosting Regressor: -**

```
=
[185] Python
... MAE: 1.705947689373957
    MSE: 10.90467045731623
    R2 score: 0.8947487714525375
    RMSE : 3.302222
```


- **XGB Regressor: -**

```
... MAE: 1.368200062168464
MSE: 9.43618017020281
R2 score: 0.9115482386132322
RMSE : 3.071837
```

Lower values of RMSE indicate better fit. The lower the RMSE, the better a given model is able to "fit" a dataset. here in XG BOOST regressor it implies highest R2score and lowest RMSE among all other models

- **Hyper Parameter Tuning the Best Score: -**

To Display the Best Score

```
gs_results.best_score_
[92] Python
0.9071962570393127
```

- **Cross Validation: -**

```
from sklearn.model_selection import KFold, cross_val_score, train_test_split
random_model = RandomForestRegressor(n_estimators=300, max_depth=4)
#Fit
random_model.fit(X_train, Y_train)
xgb_accuracy = round(random_model.score(X_train, Y_train)*100, 2)
print(round(xgb_accuracy, 2), '%')
[209] Python
... 84.26 %
```

- Here we have concluded that after analyzing the parameter and the matrices of all the models we have used that XG Boost Regressor is the best Model among all.
- As we have applied the hyper parameter tuning and derived the best parameter with the best score of 90%
- Also, we have done the Cross Validation Of the model to Cross verify the Models and we got XG Boost regressor as the highest scoring model when compared with other

Conclusion and Future works

The increased prices of new cars and the financial incapability of the customers to buy them, Used Car sales are on a global increase. Therefore, there is an urgent need for a Used Car Price Prediction system which effectively determines the worthiness of the car using a variety of features. The proposed system will help to determine the accurate price of used car price prediction. This paper compares 3 different algorithms for machine learning: Linear Regression, random forest regression and XG Boost Regression.

In future this machine learning model may bind with various website which can provide real time data for price prediction. Also, we may add large historical data of car price which can help to improve accuracy of the machine learning model. We can build an android app as user interface for interacting with user. For better performance, we plan to judiciously design deep learning network structures, use adaptive learning rates and train on clusters of data rather than the whole dataset.

References

- <https://www.ijraset.com/research-paper/price-prediction-of-used-cars-using-machine-learning>
- [https://www.researchgate.net/publication/343878698 Used Cars Price Prediction using Supervised Learning Techniques](https://www.researchgate.net/publication/343878698_Used_Cars_Price_Prediction_using_Supervised_Learning_Techniques)
- <https://www.irjet.net/archives/V8/i4/IRJET-V8I4278.pdf>
- <https://towardsdatascience.com/used-car-price-prediction-using-machine-learning-e3be02d977b2>